

Academic Skills in Computer Science

Maxim Borisyak, Andrey Ustyuzhanin

Constructor University Bremen

February 10, 2026

The Scientific Process

What is Science?

Science: systematic understanding of the natural world through observation, experimentation and reasoning.

Develops theories that explain phenomena and make predictions.

The Scientific Method

An iterative process:

1. **Observation:** notice phenomena in the world;
2. **Question:** formulate questions about what you observed;
3. **Hypothesis:** propose a testable explanation;
4. **Prediction:** derive testable predictions from the hypothesis;
5. **Experimentation:** design and conduct experiments to test predictions;
6. **Analysis:** analyze data and draw conclusions;
7. **Iteration:** refine, alter, or reject the hypothesis based on results.

Falsifiability (Karl Popper)

A scientific hypothesis must be falsifiable:

- there must exist a possible observation that would disprove it;
- example: "All swans are white" is falsifiable
 - finding a black swan disproves it;
- non-falsifiable statements are not scientific.

Reproducibility

Independent researchers must be able to reproduce experiments:

- results must be consistent across different settings;
- validates findings;
- fundamental to scientific method.

Parsimony (Occam's Razor)

Prefer simpler explanations:

- simpler hypotheses over complex ones;
- simpler explanations have fewer assumptions;
- complex explanations require stronger evidence.

Mathematics vs. Science

Mathematics: The Language of Patterns

Nature of Mathematics:

- **deductive reasoning**: starts from axioms and derives theorems;
- **proof-based**: mathematical statements are proved, not tested;
- **absolute truth** (within a formal system);
- **abstract**: not necessarily tied to physical reality;
- **timeless**: mathematical truths don't change with new observations.

Mathematical Process

Mathematical process:

1. define axioms and rules;
2. propose conjectures;
3. construct proofs or find counterexamples;
4. establish theorems.

Example: Pythagorean theorem ($a^2 + b^2 = c^2$) is proven from Euclidean axioms, not tested experimentally.

Science: Understanding Nature

Nature of Science:

- **inductive reasoning**: generalizes from observations;
- **evidence-based**: theories supported by empirical evidence;
- **provisional truth**: theories are best current explanations;
- **empirical**: tied to observable reality;
- **evolving**: theories are refined or replaced with new evidence.

Scientific Process

Scientific process:

1. observe phenomena;
2. formulate hypotheses;
3. test through experiments;
4. develop theories.

Example: Newton's laws were refined by Einstein's relativity when new evidence emerged (high velocities, strong gravity).

Key Differences: Mathematics vs. Science

Aspect	Mathematics	Science
Method	Deductive (axioms → theorems)	Inductive (observations → theories)
Validation	Proof	Experimental evidence
Truth	Absolute (in formal system)	Provisional
Change	Theorems remain valid in their system	Theories evolve with new evidence
Subject	Abstract structures	Natural world

The Relationship Between Math and Science

Mathematics and science are interconnected:

- **mathematics as a tool**: science uses mathematics to model phenomena;
- **science as inspiration**: physical problems inspire mathematical development;
- **applied mathematics**: develops mathematical methods for scientific problems;
- **mathematical physics**: studies mathematical structures underlying physical theories.

Computer Science: The Science of Computation

What is Computer Science?

The study of computation, information and algorithms.

Not programming! Fundamental questions:

- what can be computed? (computability);
- how efficiently? (complexity);
- what are the limits of computation?;
- how do algorithms behave?

Mathematical foundations:

- **computability theory**: what problems are solvable?
 - Turing machines, decidability, halting problem;
- **complexity theory**: resources needed for computation
 - P vs NP, time/space complexity classes;
- **formal verification**: proving program correctness;
- **algorithm analysis**: proving runtime and space bounds.

Scientific study of computational phenomena:

- **algorithm behavior:** measuring performance in practice
 - average-case vs. worst-case;
 - observed performance vs. theoretical bounds;
- **systems evaluation:** measuring system performance
 - caching strategies, database indexing;
 - network protocols, distributed systems;
- **experimental algorithmics:** testing algorithmic hypotheses.

Machine Learning as Computer Science

Machine learning: algorithms that learn from data.

Theoretical aspects:

- PAC learning theory, VC dimension;
- optimization theory (gradient descent, convex optimization);
- statistical learning theory (bias-variance tradeoff);
- information theory (entropy, mutual information).

Machine Learning: Empirical Study

Scientific methods in ML:

- hypothesis: model architecture will capture patterns;
- experiment: train on data, validate on held-out set;
- analysis: measure generalization (accuracy, F1, AUC);
- iteration: refine based on empirical results.

Open questions:

- why do deep networks generalize?
- how much data is needed?
- what representations do models learn?

Types of Knowledge in CS

Computer science produces different types of knowledge: **Proven facts** (mathematical):

- sorting requires $\Omega(n \log n)$ comparisons;
- halting problem is undecidable;
- binary search is $O(\log n)$.

Empirical findings (scientific):

- hash tables perform well in practice;
- neural networks learn hierarchical features;
- cache locality affects performance.

Research Methods in Computer Science

CS research methods:

1. **mathematical/theoretical**: proving theorems
 - complexity lower bounds, correctness proofs;
2. **experimental**: measuring algorithmic performance
 - benchmark datasets, performance evaluation;
3. **empirical**: studying computational phenomena
 - user studies, large-scale data analysis;
4. **simulation**: modeling system behavior.

Structure of Scientific Papers

Scientific Papers (Empirical Sciences)

Standard structure (IMRaD):

1. **introduction**: motivate the problem, state research question;
2. **methods**: describe experimental setup and procedures;
3. **results**: present data and observations;
4. **discussion**: interpret results, compare with prior work;
5. **conclusion**: summarize findings and implications.

Scientific Papers: Key Characteristics

Empirical focus:

- report experimental data;
- describe methodology in detail (reproducibility);
- statistical analysis of results;
- figures and tables with data.

Citation style:

- cite prior experimental work;
- compare results with literature;
- discuss discrepancies.

Mathematical Papers

Standard structure:

1. **introduction**: motivate the problem, state main theorems;
2. **preliminaries**: definitions, notation, known results;
3. **main results**: theorems with proofs;
4. **corollaries and applications**: consequences of main results;
5. **conclusion**: summarize contributions, open problems.

Mathematical Papers: Key Characteristics

Proof-based:

- every claim must be proven;
- rigorous logical structure;
- definitions must be precise;
- no experimental data.

Theorem-proof format:

- state theorem formally;
- provide complete proof;
- lemmas support main theorems.

Computer Science Papers: Theoretical

Similar to mathematical papers:

1. **introduction**: problem, main results;
2. **preliminaries**: definitions, model, prior results;
3. **main results**: theorems and proofs;
4. **conclusion**: contributions, open problems.

Examples: complexity lower bounds, algorithm correctness proofs, impossibility results.

Computer Science Papers: Empirical

Similar to scientific papers:

1. **introduction**: problem, research questions;
2. **related work**: prior approaches;
3. **method/system**: algorithm or system description;
4. **experimental setup**: datasets, baselines, metrics;
5. **results**: performance measurements, analysis;
6. **conclusion**: findings, future work.

Examples: ML papers, systems papers, performance studies.

Computer Science Papers: Hybrid

Combine both approaches:

1. **introduction**: problem, contributions;
2. **related work**: theoretical and empirical background;
3. **theoretical analysis**: algorithm with proofs;
4. **experimental evaluation**: empirical validation;
5. **conclusion**: theoretical and empirical contributions.

Example: new algorithm with proven bounds AND experimental comparison to existing methods.

Example: Auto-Encoding Variational Bayes

Auto-Encoding Variational Bayes

Diederik P. Kingma

Machine Learning Group
Universiteit van Amsterdam
dpkingma@gmail.com

Max Welling

Machine Learning Group
Universiteit van Amsterdam
welling.max@gmail.com

Abstract

How can we perform efficient inference and learning in directed probabilistic models, in the presence of continuous latent variables with intractable posterior distributions, and large datasets? We introduce a stochastic variational inference and learning algorithm that scales to large datasets and, under some mild differentiability conditions, even works in the intractable case. Our contributions are two-fold. First, we show that a reparameterization of the variational lower bound yields a lower bound estimator that can be straightforwardly optimized using standard stochastic gradient methods. Second, we show that for i.i.d. datasets with continuous latent variables per datapoint, posterior inference can be made especially efficient by fitting an approximate inference model (also called a recognition model) to the intractable posterior using the proposed lower bound estimator. Theoretical advantages are reflected in experimental results.

1 Introduction

Example: Auto-Encoding Variational Bayes

1 Introduction

How can we perform efficient approximate inference and learning with directed probabilistic models whose continuous latent variables and/or parameters have intractable posterior distributions? The variational Bayesian (VB) approach involves the optimization of an approximation to the intractable posterior. Unfortunately, the common mean-field approach requires analytical solutions of expectations w.r.t. the approximate posterior, which are also intractable in the general case. We show how a reparameterization of the variational lower bound yields a simple differentiable unbiased estimator of the lower bound; this SGVB (Stochastic Gradient Variational Bayes) estimator can be used for efficient approximate posterior inference in almost any model with continuous latent variables and/or parameters, and is straightforward to optimize using standard stochastic gradient ascent techniques.

For the case of an i.i.d. dataset and continuous latent variables per datapoint, we propose the Auto-Encoding VB (AEVB) algorithm. In the AEVB algorithm we make inference and learning especially efficient by using the SGVB estimator to optimize a recognition model that allows us to perform very efficient approximate posterior inference using simple ancestral sampling, which in turn allows us to efficiently learn the model parameters, without the need of expensive iterative inference schemes (such as MCMC) per datapoint. The learned approximate posterior inference model can also be used for a host of tasks such as recognition, denoising, representation and visualization purposes. When a neural network is used for the recognition model, we arrive at the *variational auto-encoder*.

Example: Auto-Encoding Variational Bayes

efficiently learn the model parameters, without the need of expensive iterative inference schemes (such as MCMC) per datapoint. The learned approximate posterior inference model can also be used for a host of tasks such as recognition, denoising, representation and visualization purposes. When a neural network is used for the recognition model, we arrive at the *variational auto-encoder*.

2 Method

The strategy in this section can be used to derive a lower bound estimator (a stochastic objective function) for a variety of directed graphical models with continuous latent variables. We will restrict ourselves here to the common case where we have an i.i.d. dataset with latent variables per datapoint, and where we like to perform maximum likelihood (ML) or maximum a posteriori (MAP) inference on the (global) parameters, and variational inference on the latent variables. It is, for example,

Example: Auto-Encoding Variational Bayes

and stochastic variational inference using the reparameterization trick we describe in this paper. Their work was developed independently of ours and provides an additional perspective on AEVB.

5 Experiments

We trained generative models of images from the MNIST and Frey Face datasets³ and compared learning algorithms in terms of the variational lower bound, and the estimated marginal likelihood.

The generative model (encoder) and variational approximation (decoder) from section 3 were used, where the described encoder and decoder have an equal number of hidden units. Since the Frey Face data are continuous, we used a decoder with Gaussian outputs, identical to the encoder, except that the means were constrained to the interval (0, 1) using a sigmoidal activation function at the

³Available at <http://www.cs.nyu.edu/~roweis/data.html>

Example: Auto-Encoding Variational Bayes

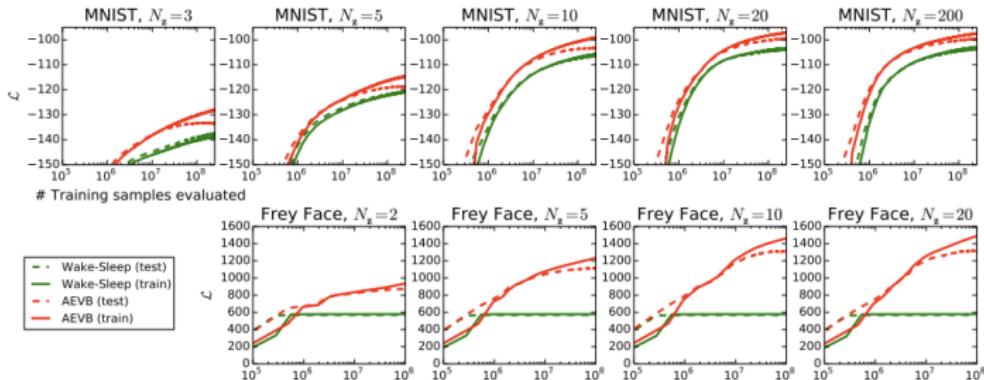


Figure 2: Comparison of our AEVB method to the wake-sleep algorithm, in terms of optimizing the lower bound, for different dimensionality of latent space (N_z). Our method converged considerably faster and reached a better solution in all experiments. Interestingly enough, more latent variables does not result in more overfitting, which is explained by the regularizing effect of the lower bound. Vertical axis: the estimated average variational lower bound per datapoint. The estimator variance was small (< 1) and omitted. Horizontal axis: amount of training points evaluated. Computation took around 20-40 minutes per million training samples with a Intel Xeon CPU running at an effective 40 GFLOPS.

Example: Auto-Encoding Variational Bayes

6 Conclusion

We have introduced a novel estimator of the variational lower bound, Stochastic Gradient VB (SGVB), for efficient approximate inference with continuous latent variables. The proposed estimator can be straightforwardly differentiated and optimized using standard stochastic gradient methods. For the case of i.i.d. datasets and continuous latent variables per datapoint we introduce an efficient algorithm for efficient inference and learning, Auto-Encoding VB (AEVB), that learns an approximate inference model using the SGVB estimator. The theoretical advantages are reflected in experimental results.

7 Future work

Since the SGVB estimator and the AEVB algorithm can be applied to almost any inference and learning problem with continuous latent variables, there are plenty of future directions: (i) learning hierarchical generative architectures with deep neural networks (e.g. convolutional networks) used for the encoders and decoders, trained jointly with AEVB; (ii) time-series models (i.e. dynamic Bayesian networks); (iii) application of SGVB to the global parameters; (iv) supervised models with latent variables, useful for learning complicated noise distributions.

Paper Structure Reflects Methodology

Paper structure follows the discipline's methodology.

Science papers follow the scientific method:

- introduction: research question, hypothesis;
- methods: how to test the hypothesis;
- results: data from experiments;
- discussion: interpret results, compare to hypothesis;
- mirrors: hypothesis → experiment → analyze → conclude.

Paper Structure Reflects Methodology

Mathematics papers follow deductive reasoning:

- introduction: state theorems to be proven;
- preliminaries: axioms, definitions, prior results;
- main results: prove theorems from axioms;
- conclusion: what was proven, open problems;
- mirrors: axioms → conjecture → proof → theorem.

Paper Structure Reflects Methodology

CS papers follow their methodology:

- theoretical: like mathematics (axioms → proof);
- empirical: like science (hypothesis → experiment);
- hybrid: both approaches in one paper;
- structure matches the type of claims being made.

Peer Review Process

Papers are evaluated before publication:

- authors submit manuscript to journal or conference;
- editor sends to multiple expert reviewers;
- reviewers evaluate correctness, novelty, significance;
- accept, reject, or request revisions;
- iterative process until acceptance or rejection.

Peer Review: Purpose

Why peer review:

- finds errors in proofs or experiments;
- evaluates significance of contributions;
- ensures paper meets publication standards;
- improves quality through reviewer feedback;
- science is self-correcting.

Conclusion

Summary

Different disciplines, different methods:

- **science**: inductive reasoning, experiments, provisional truth;
- **mathematics**: deductive reasoning, proofs, absolute truth;
- **computer science**: combines both approaches.

Paper structure reflects methodology.

Understand the methodology to succeed in research.